

# Intro to R Programming

-- *Nikhil Vidhani*

# What is a program?

- Translation of an algorithm into a language that computer understands
- An algorithm takes input, perform some operations and gives output
  - Executes in finite time
  - E.g. sorting, searching, reading, copying!
- Complexity of a Program
  - Time and space!
  - E.g. Fibonacci series!
- Programming Paradigms
  - Iterative vs Recursive
  - Procedural vs Object Oriented
- Good Program
  - Re-readable, organized and modular

# Typical Programing Errors

- Syntactical (spelling mistake)
  - Will get caught very easily! Just run the program.
- Semantic Errors (meaningless operations)
  - For e.g. “nikhil”+32
  - Exceptions: like divide by 0.
    - Note that R catches this exception and returns Inf
  - May get caught. A warning will be thrown nonetheless.
- Logical Errors (Unintentional)
  - Program will crash, run forever or give a wrong answer!
  - Debugging requires some skill and experience.

# Getting Help in R

- From Console

- Just type: ? followed by function name without parenthesis
- E.g. `?mean`; `?sum`; `?length`;
- Clarify:
  - `?mean` - help for the function “mean” in the available libraries
    - Same as `help(mean)`
  - `??mean` - will perform a keyword search over the entire database
    - Same as `help.search("mean")`
  - `mean()` - call the function mean
  - `mean` - print the definition of the function “mean”

- From Web sources

- Most reliable and easy to incorporate is [www.stackoverflow.com](http://www.stackoverflow.com).
- [www.r-bloggers.com](http://www.r-bloggers.com) is also quite helpful.
- You can use <https://cran.r-project.org> for any resource on R
- Even typing your question in google will get you good results!
  - 99% of your questions are already answered! You just need to find them!

# R Input and Output

- Simple assignment
  - `X = 1;` (or `X <- 1;`)
  - Assignment is always right to left
    - Read 1 goes into X
    - We aren't comparing X with 1 here
  - The semi-colon isn't necessary in R, but it's a good practice to use it
  - `X = ;` is incomplete
  - `#` (prefix) is used as a comment. Use it for helpful comments.
  - Use Ctrl-Shift-C for multi-line comments
- Value of X can be seen by
  - `X;`

# Vectors

- A sequence of numbers. Many ways to input!
  - `Y = c(1,7,-3,41);` # concatenate arbitrary numbers
  - `Y = 1:10;` # natural numbers
  - `Y = seq(1,100,9);` # skip by 9
  - `Y = rep(2, 3);` # repeat 3 times
  - `Y = rep(1:2, 3);` # repeat the vector
  - `Y = rep(1:2, each = 3);` # repeat each element 3 times
  - `Y = c();` # empty vector
  - Execute this: `c(1:3, rep(c(5,7), each = 2), rep(9, 4), 7);`
- Length of vector: `length(Y);`
- Accessing  $i^{\text{th}}$  element of vector: `Y[i];` # square brackets
  - `i` should be between 1 and `length(Y)`
  - Printing the entire vector is as before: `Y;`

# Objects in R

- 5 basic (atomic) types of objects
  - character – strings
  - numeric – real numbers. Also called double.
  - integer – natural numbers. Default data type for numeric vectors.
    - `typeof(1:10)`
  - complex – complex numbers. We won't use them now!
  - logical – True/False (binary)
- Most basic collection of objects is a vector (also called an array)
  - Can only contain objects of same class (i.e. character or integer; not both)
  - “list” is a special type of object and can contain heterogeneous objects
    - Any Combination of vector, matrix, atomic types etc.
    - It can even contain another list as an object. E.g. linked-lists!
    - Due to its generality its very slow and hence rarely used with large datasets unless situation demands it

# Numbers

- Default type of any number is numeric (i.e. real). `typeof(1)`
- R can differentiate between corner cases:
  - `1/0` is `Inf` -- `is.infinite()`;
  - `0/0` is `NaN` -- `is.nan()`;
  - Missing data is `NA` -- `is.na()`;
  - Check what's `Inf-Inf` ?
- Arithmetic Operations
  - `*` multiplies
  - `/` divides
  - `^` takes exponent
  - `%%` is the modulo (remainder) operator. Try: `7 %% 2`;



# Coercion

- Mixing Objects

- Automatically coerced to the same class.
- Try: `c(1:7, "a"); c(T, 2); c("a", FALSE);`
- Implicit coercion!
- Never use unless you know what you're doing!

- Explicit Coercion

- `as.character(1:5);`
- `as.numeric("iimb"); # warning!`
- `as.logical(seq(-2,2,1));`

# List

- Can carry different types of data together
  - `L = list(1, FALSE, 3.14, "iimb", "c", 4-3i);`
  - Print list: `L;`
  - L is in fact a list of lists. Check: `typeof(L); typeof(L[4]); typeof(L[[4]]);`
  - Single square brackets `[i]` access the  $i^{\text{th}}$  list embedded in the list `L`
  - Double square brackets `[[i]]` access the  $i^{\text{th}}$  element
  - Can append elements in list: `L = append(L, "7th");`
  - `unlist(L);` will coerce all elements into a single type and return a vector
  - Delete an element from a list:
    - I don't know how to do that!
    - Let's google: "delete element from list in R"
    - Open the answer on [www.stackoverflow.com](http://www.stackoverflow.com)

# Matrices

- Generalization of vectors
  - 2 dimensions instead on one!
  - $N \times K$  matrix means a matrix having  $N$  rows and  $K$  columns. Total of  $NK$  elements.
  - `M = matrix(nrow = 2, ncol = 3);`
  - Dimensions: `dim(M);`
  - Can think of  $M$  as
    - 3 columns vectors each of length 2, *or*
    - 2 row vectors each of length 3
  - Populate matrix: `M = rbind(1:3, 4:6);`
  - Alternatively populate as: `M = cbind(1:2, 3:4, 5:6);`

# Matrices (cont.)

- Indexing a matrix
  - `M[i,j]` gives the element at  $i^{\text{th}}$  row and  $j^{\text{th}}$  column
  - `M[i, ]` gives the entire  $i^{\text{th}}$  row (a vector)
  - `M[,j]` gives the entire  $j^{\text{th}}$  column (a vector)
- Matrix multiplication
  - `*` just does an element wise multiplication, i.e.  $(M * N)_{ij} = M_{ij} * N_{ij}$
  - `%*%` performs the usual matrix multiplication. Try: `M %*% N`
    - Dimensions must match
    - Try `t(M) %*% M;`
    - `t(M)` takes transpose of a matrix!

# Matrices (cont.)

- Identity matrix: `diag(3)`
- Diagonal Matrix: `diag(c(1,5,7)); diag(1:7);`
- Diagonal of a matrix: `diag(M)`
- Trace of a matrix: `sum(diag(M))`
- Inverse of a matrix:
  - Must be a square matrix: `M = matrix(1:9, nrow = 3, ncol = 3);`
    - Another way to create a matrix. Data is entered column-wise.
  - Determinant must be non-zero: `det(M); M[3,3] = 19; det(M);`
  - Inverse: `solve(M);`

# Reading Data

- Download some stock data from NSE
  - [https://www.nseindia.com/products/content/equities/indices/historical\\_index\\_data.htm](https://www.nseindia.com/products/content/equities/indices/historical_index_data.htm)
  - Save the CSV file as data.csv
- From CSV (most common)
  - `setwd("C:/Users/nikhil/data.csv"); nifty = read.csv("data.csv");`
  - Alternatively: `nifty = read.csv(" C:/Users/nikhil/data.csv");`
- From Excel
  - Search it yourself! It is not recommended btw.
- From clipboard
  - `read.table("clipboard");`
  - This is quick fix for small data transfer between R and excel. Use `read.csv()` as your primary method for data reading!

# Reading Data (cont.)

- Viewing data

- `View(nifty);`

- Date

- `nifty$Date = as.Date(nifty$Date, format = "%d-%b-%Y");`

- `n = nrow(nifty);`

- `d = nifty$Date[1];`

- `format(d, format = "%D");` # 04/02/18

- `format(d, format = "%d-%m-%y");` # 02-04-18

- `format(d, format = "%d.%b.%Y");` # 02.Apr.2018

- `format(d, format = "%d_%B_%Y");` # 02\_April\_2018

- Alternatively,

- `read.table("data.csv", header = T, sep = ",", nrow = 5);`

# if-else

- `if(<COND_1>) {  
 # do something!  
}`
- `if(<COND_1>) {  
 # do something!  
} else {  
 # ...  
}`
- `if(<COND_1>) {  
 # do something!  
} else if(<COND_2>) {  
 # ...  
} else {  
 # ...  
}`

```
if(nifty$Close[2] > nifty$Close[1]) {  
  str = paste("Stock market closed green on", nifty$Date[2]);  
} else if(nifty$Close[2] > nifty$Open[2]) {  
  str = paste("Stock market closed above opening on", nifty$Date[2]);  
} else {  
  str = paste("Stock market was red and closed below opening on", nifty$Date[2]);  
}  
print(str);
```



# for loop

- Looping is used to perform similar set of tasks repetitively
  - `for(i in n:1) {  
 print(nifty$Date[i]);  
}`
  - `n:1`; is same as `seq(n,1,1)`; i.e. backwards counting!
  - Alternatively, you can execute: `rev(nifty$Date)`; or `nifty$Date[n:1]`;
- Try avoiding loops if you can!
  - Increasing all dates by a week: `nifty$Date + 7`
  - Finding Daily growth: `nifty$Close[-1] / nifty$Close[-n]`
  - Daily diff. b/w high and low prices: `nifty$High - nifty$Low`
  - Question: find % growth in daily volatility
    - Volatility is defined as:  $Vol_t = (High_t - Low_t) / Open_t$
    - Percentage Growth is defined as:  $\%G = \frac{(Value_{t+1} - Value_t)}{Value_t} * 100$
- Loops and if-else can be used in conjunction!

# Jumping

- Till now all our commands executed sequentially
- There may be circumstances when we need to jump
- Next and Break
  - `next` is used to skip an iteration, while `break` exits the loop entirely.
  - ```
for(i in 1:10) {  
  if(i <= 3) {  
    next;  
  }  
  if(i > 6) {  
    break;  
  }  
  print(i);  
}  
i;
```
  - `return()` is used to exit a function with a value.

# Function

- Organize often-used set of instructions separately in a “function”
- Calling a function will execute all the commands in the body of function
- We have used many functions till now
  - They end with parenthesis: `()`
  - E.g. `sum()`; `rbind()`; `vector()`; `format()`; `read.csv()`; etc
  - Note that curly braces `{}` are used for if-else and for loops, square braces `[]` for vector/matrix indexing and parenthesis `()` for grouping, if-else & for conditions and functions.
- A function has
  - A name by which we call them, e.g. `sum`
  - A set of inputs to be put within parenthesis like numbers `1:10` in `sum()`
  - Return value which is the output of the function like the sum of numbers in `sum()`

# Function Example

```
my_mean = function(x) {  
  n = length(x);  
  mean = sum(x) / n;  
  return(mean);  
}
```

- Name of the function is: `my_mean`
- Input is: `x`
- Output is: `mean`
  - Note that the mean here is just a name, we could well have used any other name without changing anything about our function

# Function (Example) Cont.

- Alternate ways to write the same function

- ```
my_mean = function(x) {  
  return( sum(x) / length(x) );  
}
```

- No need to store sum and length. We can directly divide them!

- ```
my_mean = function(x) {  
  sum(x) / length(x);  
}
```

- No need for an explicit return. The last statement is returned by default.

- Try various value with `my_mean()` and the inbuilt `mean()`. See that the answers are exactly the same.

- Write a function for variance where  $Var(x) = mean([x - mean(x)]^2)$ 
  - Compare it with the inbuilt `var()` function?

# Multiple conditions & which() function

- The arguments to `if()` and `which()` and the output of `is.xx()` family of functions is a logical object, i.e. either `TRUE` or `FALSE`.
- A valid combination of logical objects is also a logical object. E.g.
  - Logical AND: `TRUE & FALSE` is `FALSE`
  - Logical OR: `TRUE | FALSE` is `TRUE`
  - Logical NOT: `! FALSE` is `TRUE`
- De Morgan's Law
  - $!(A | B) = (!A) \& (!B)$
  - $!(A \& B) = (!A) | (!B)$

- The below two indexes are one and same (by De Morgan Law),
  - `day = as.numeric( substr(nifty$Date, 9, 10) );`
    - OR `day = as.numeric( format(df[,1], format = "%d") );`
  - `idx = which(nifty$Close > nifty$Open & (day < 5));`
  - `idx = which( !(nifty$Close < nifty$Open | (day >= 5)) );`
- `which()` gives the indexes matching the criterion. E.g. out of `101:200` which numbers are multiples of 2,3 and 5 ?
  - `count = 101:200;`
  - `which( count %% 2 == 0 & count %% 3 == 0 & count %% 5 == 0 );`
  - `count[count %% 2 == 0 & count %% 3 == 0 & count %% 5 == 0];`
- We can do multi-way match using `%in%`
  - `mult_17 = seq(17,300,17);`
  - `which(count %in% mult_17);`
  - `which(mult_17 %in% count);`
  - `which(!(count %in% mult_17));`
  - `which(!(mult_17 %in% count));`

# Correlation

- Difference between Independence and Correlation
  - $X \sim N(0,1)$ ,  $Y = X^2$ . Are  $X$  and  $Y$  correlated?
  - Let's check in R
    - ```
X = rnorm(1000, 0, 1);  
Y = X^2;  
cor(X,Y);
```
    - However, 


```
cor(X,X^3) != 0.
```
  - Let  $Y \sim N(0,1)$  independent of  $X$ . Then is  $\text{cor}(X,Y) = 0$ ? What about  $\text{cor}(X^q, Y)$ ?
    - ```
Y = rnorm(1000,0,1);
```
    - ```
cor_q = rep(NA, 10);  
for(q in 1:10) {  
    cor_q[q] = cor(X^q, Y);  
}  
cor_q;
```
- Independence implies NO correlation of any functional form.




# Regression


- Let,  $Y = \beta_0 + \beta_1 \cdot X + u$  be the true model. By regressing  $Y$  on  $X$ , we hope to recover an unbiased estimate of  $\beta_1$  and see how much of the variation in  $Y$  is explained by variation in  $X$  unrelated to variation in  $u$ .

```
n = 1000;  
x = rnorm(n, 0, 1);  
u = rnorm(n, 0, 0.1);  
y = u + x;  
plot(x,y);
```




```
fit = lm(y ~ x);  
summary(fit);  
stargazer(fit, type = "html", out = "fit.html");  
# Regression Line  
abline(fit$coefficients, col = "red", lwd = 2);
```





```
u = rnorm(n, 0, 0.5);  
y = u + x;  
plot(x,y);  
fit = lm(y ~ x);  
summary(fit);  
# Regression Line  
abline(fit$coefficients, col = "red", lwd = 2);
```

```
u = rnorm(n, 0, 5);  
y = u + x;  
plot(x,y);  
fit = lm(y ~ x);  
summary(fit);  
# Regression Line  
abline(fit$coefficients, col = "red", lwd = 2);
```



# Some Useful functions

- `unique()`, `duplicated()`
- `list.files()`
  - Pattern matching using regex
  - All files starting from "s": `"^s"`
  - All files starting with "b" or "d": `"^(b|d)"`
  - All CSV files: `".*.csv"`
- `order()`
  - Sort data/dataframes
  - It gives the sequence of ordered indexes NOT the ordered numbers
  - Can do 2-way and 3-way sorts
- `union()`, `intersect()`
- `cumsum()`, `cumprod()`
  - `cumprod()` using only `cumsum()` ?

# Merging

- Combining two datasets is a very routine and important task
- Merging is akin to Joining (in relational database, SQL etc)

□ Summary:

$r =$

attr1	attr2
a	r1
b	r2
c	r3

$s =$

attr1	attr3
b	s2
c	s3
d	s4

$r \bowtie s$

attr1	attr2	attr3
b	r2	s2
c	r3	s3

$r \bowtie_s s$

attr1	attr2	attr3
a	r1	<i>null</i>
b	r2	s2
c	r3	s3

$r \bowtie_r s$

attr1	attr2	attr3
b	r2	s2
c	r3	s3
d	<i>null</i>	s4

$r \bowtie_{rs} s$

attr1	attr2	attr3
a	r1	<i>null</i>
b	r2	s2
c	r3	s3
d	<i>null</i>	s4

# Merging Example

- Taken from [Stackoverflow webpage](#)
- Create two datasets:
  - `df1 = data.frame(cust_id = c(1:6), Product = c(rep("Toaster", 3), rep("Radio", 3)));`
  - `df2 = data.frame(cust_id = c(2, 4, 6), State = c(rep("Alabama", 2), rep("Ohio", 1)));`
  - The data looks like:
- We will need to use a new package `dplyr`
  - `install.packages("dplyr");`

<u>df1</u>		<u>df2</u>	
<u>cust_id</u>	<u>Product</u>	<u>cust_id</u>	<u>State</u>
1	Toaster	2	Alabama
2	Toaster	4	Alabama
3	Toaster	6	Ohio
4	Radio		
5	Radio		
6	Radio		

# Inner Join

- `dplyr::inner_join(df1, df2);`

<u>cust_id</u>	<u>Product</u>	<u>State</u>
2	Toaster	Alabama
4	Radio	Alabama
6	Radio	Ohio

# Full Join

- `dplyr::full_join(df1, df2);`

<u>cust_id</u>	<u>Product</u>	<u>State</u>
1	Toaster	NA
2	Toaster	Alabama
3	Toaster	NA
4	Radio	Alabama
5	Radio	NA
6	Radio	Ohio



# Left Join

- `dplyr::left_join(df1, df2);`

<u>cust_id</u>	<u>Product</u>	<u>State</u>
1	Toaster	NA
2	Toaster	Alabama
3	Toaster	NA
4	Radio	Alabama
5	Radio	NA
6	Radio	Ohio

# Right Join

- `dplyr::right_join(df1, df2);`

- The above is identical to,  
`dplyr::left_join(df2, df1);`

<u>cust_id</u>	<u>Product</u>	<u>State</u>
2	Toaster	Alabama
4	Radio	Alabama
6	Radio	Ohio

# Cartesian Product

- Every row of `df1` multiplied with every row of `df2`
- `cart_prd = merge(df1, df2, by = NULL);`

<u>S.No.</u>	<u>cust id.x</u>	<u>Product</u>	<u>cust id.y</u>	<u>State</u>
1	1	Toaster	2	Alabama
2	2	Toaster	2	Alabama
3	3	Toaster	2	Alabama
4	4	Radio	2	Alabama
5	5	Radio	2	Alabama
6	6	Radio	2	Alabama
7	1	Toaster	4	Alabama
8	2	Toaster	4	Alabama
9	3	Toaster	4	Alabama

<u>S.No.</u>	<u>cust id.x</u>	<u>Product</u>	<u>cust id.y</u>	<u>State</u>
10	4	Radio	4	Alabama
11	5	Radio	4	Alabama
12	6	Radio	4	Alabama
13	1	Toaster	6	Ohio
14	2	Toaster	6	Ohio
15	3	Toaster	6	Ohio
16	4	Radio	6	Ohio
17	5	Radio	6	Ohio
18	6	Radio	6	Ohio

- Cartesian products are extremely slow. Never do that even on a decent sized (> 1000 rows) dataset. Your computer will probably hang.
- Although there is no use of Cartesian products, it encapsulates all types of merges. Meaning we can extract any type of merge from a Cartesian product.
- Inner join can be extracted via
  - `idx = which(cart_prd$cust_id.x == cart_prd$cust_id.y);`
  - `cart_prd[idx,];`

# Merging with more than one variable

- Most of the merging usage is with two variables: date and company name. Generate data using below:

```
date = seq.Date(as.Date("2018-01-01"), by = 1, length.out = 5);
comp = c("A", "B", "C", "D", "E");
all_pairs = merge(comp, date, by = NULL);
# sales data - 15 points
idx = sample(1:nrow(all_pairs), 15, replace = F);
df1 = data.frame(comp = all_pairs$x[idx], date = all_pairs$y[idx],
                 sales = round(runif(15, min = 1e3, max = 1e5)));
# advertising data - 12 points
idx = sample(1:nrow(all_pairs), 12, replace = F);
df2 = data.frame(comp = all_pairs$x[idx], date = all_pairs$y[idx],
                 adv = round(runif(12, min = 1e2, max = 1e4)));
```

<u>df1</u>				<u>df2</u>		
<u>comp</u>	<u>date</u>	<u>sales</u>		<u>comp</u>	<u>date</u>	<u>adv</u>
B	05-04-2018	53429		C	03-04-2018	980
A	03-04-2018	70750		C	06-04-2018	6183
B	03-04-2018	5426		B	06-04-2018	8077
B	06-04-2018	88504		B	02-04-2018	241
E	06-04-2018	78449		D	02-04-2018	3642
A	04-04-2018	98954		B	04-04-2018	3878
D	04-04-2018	57618		E	02-04-2018	1501
E	04-04-2018	22011		B	03-04-2018	6727
C	02-04-2018	85976		D	06-04-2018	4259
D	02-04-2018	42842		D	03-04-2018	4434
E	02-04-2018	94057		A	06-04-2018	6745
A	05-04-2018	25063		E	03-04-2018	2208
E	05-04-2018	51320				
D	06-04-2018	99720				
E	03-04-2018	16845				

# Inner Join

- `dplyr::inner_join(df1, df2);`

<u>comp</u>	<u>date</u>	<u>sales</u>	<u>adv</u>
B	03-04-2018	5426	6727
B	06-04-2018	88504	8077
D	02-04-2018	42842	3642
E	02-04-2018	94057	1501
D	06-04-2018	99720	4259
E	03-04-2018	16845	2208

# Full Join

- `dplyr::full_join(df1, df2);`

<u>comp</u>	<u>date</u>	<u>sales</u>	<u>adv</u>
B	05-04-2018	53429	NA
A	03-04-2018	70750	NA
B	03-04-2018	5426	6727
B	06-04-2018	88504	8077
E	06-04-2018	78449	NA
A	04-04-2018	98954	NA
D	04-04-2018	57618	NA
E	04-04-2018	22011	NA
C	02-04-2018	85976	NA
D	02-04-2018	42842	3642
E	02-04-2018	94057	1501
A	05-04-2018	25063	NA
E	05-04-2018	51320	NA
D	06-04-2018	99720	4259
E	03-04-2018	16845	2208
C	03-04-2018	NA	980
C	06-04-2018	NA	6183
B	02-04-2018	NA	241
B	04-04-2018	NA	3878
D	03-04-2018	NA	4434
A	06-04-2018	NA	6745

# Left Join

- `dplyr::left_join(df1, df2);`

<u>comp</u>	<u>date</u>	<u>sales</u>	<u>adv</u>
B	05-04-2018	53429	NA
A	03-04-2018	70750	NA
B	03-04-2018	5426	6727
B	06-04-2018	88504	8077
E	06-04-2018	78449	NA
A	04-04-2018	98954	NA
D	04-04-2018	57618	NA
E	04-04-2018	22011	NA
C	02-04-2018	85976	NA
D	02-04-2018	42842	3642
E	02-04-2018	94057	1501
A	05-04-2018	25063	NA
E	05-04-2018	51320	NA
D	06-04-2018	99720	4259
E	03-04-2018	16845	2208

# Right Join

- `dplyr::right_join(df1, df2);`

- The above is identical to,  
`dplyr::left_join(df2, df1);`

<u>comp</u>	<u>date</u>	<u>sales</u>	<u>adv</u>
C	03-04-2018	NA	980
C	06-04-2018	NA	6183
B	06-04-2018	88504	8077
B	02-04-2018	NA	241
D	02-04-2018	42842	3642
B	04-04-2018	NA	3878
E	02-04-2018	94057	1501
B	03-04-2018	5426	6727
D	06-04-2018	99720	4259
D	03-04-2018	NA	4434
A	06-04-2018	NA	6745
E	03-04-2018	16845	2208

# Loop Functions

- Writing loops in a single command. Can come very handy and compact. The function name ends with “apply”.
- List of functions: `lapply()`, `apply()`, `sapply()`, `tapply()`, `mapply()`
- These functions work on a list of inputs, not just one input!
- E.g.
  - `X = list(a = 1:10, b = rnorm(100, 0, 1), c = runif(1e3, 9, 91));`
  - `lapply(X, mean); # returns a list`
  - `sapply(X, mean); # returns a vector`
- Let's say we want to find `cor(X,X^i)` for `i in 1:10` w/o writing a loop?
  - `X = rnorm(1000, 0, 1);`
  - `sapply(1:10, function(i) cor(X,X^i));`
    - Here we have used anonymous function, i.e. a function w/o a name.



- `apply()` is mostly used for applying functions on rows or cols of a matrix
  - Like taking means by rows
  - `M = matrix(1:50, nrow = 10, ncol = 5);`
  - `apply(M, 1, mean);` # mean of each row
  - `apply(M, 2, mean);` # mean of each col
  - You can do the above using a loop also, but `apply()` is more compact.
  - The faster version of above are also available:
    - `rowSums(x)` is equivalent to `apply(x, 1, sum)`
    - `colMeans(x)` is equivalent to `apply(x, 2, mean)`
- `tapply()` is used to apply a function over a subset of vector. Lets say,
  - `x = c(rnorm(100), runif(100));`
  - `f = rep(1:2, each = 100);`
  - `tapply(x, f, mean);`

# Pipe and the `dplyr` package

- Download the `tidyverse` package and load it: `library(tidyverse);`
  - Also load `dplyr`: `library(dplyr);`
- We will use the `gapminder` dataset.
  - Install `gapminder` package
  - `df = gapminder::gapminder; df = as.data.frame(df);`
- Selecting a particular column
  - `df %>% select(lifeExp);`
  - Performing operation on data
    - `df %>% select(lifeExp) %>% mean`
- Selecting some rows (based on a filtering criterion):
  - `df %>% filter(year > 2000)`
  - `df %>% filter(year > 2000 & country == "India")`

- Selecting some rows (based on a filtering criterion):
  - `df %>% filter(year > 2000)`
  - `df %>% filter(year > 2000 & country == "India")`
- Selecting and filtering together
  - `df %>% select(country, year, lifeExp) %>% filter(year > 2000)`
  - `df %>% select(country, year, lifeExp) %>% filter(country == "India")`
- Grouping data
  - `df %>% select(country, year, lifeExp) %>% group_by(lifeExp) %>% as.data.frame %>% head`
  - `df %>% group_by(country, gdpPercap) %>% as.data.frame %>% head`
- Create new columns (mutate)
  - `df %>% mutate(decade = 10*(as.integer(year/10))) %>% head`
- Ordering
  - `df %>% arrange(-pop) %>% head`

- Performing aggregation on grouped data is the most important usage of dplyr library
- Aggregation:
  - `df %>% select(-c(continent)) %>% group_by(country) %>% summarise_all(mean)`
  - `df %>% select(-c(continent)) %>% group_by(country) %>% summarise_all(mean)`
  - `df %>% group_by(continent) %>% summarise_all(mean)`
  - `df %>% group_by(continent, year) %>% summarise_all(mean) %>% filter(continent == "Asia")`
- Counting and ordering in Aggregation:
  - `df %>% group_by(continent, year) %>% summarise(n()) %>% filter(continent == "Asia")`
  - `df %>% group_by(continent) %>% summarise(n())`
  - `df %>% filter(year == 2007) %>% group_by(pop) %>% arrange(-pop)`
  - `df %>% filter(year == 2007) %>% group_by(pop) %>% arrange(-gdpPercap)`
  - `df %>% filter(year == 2007) %>% group_by(pop) %>% mutate(GDP = gdpPercap*pop) %>% arrange(-GDP) %>% as.data.frame %>% select(country, GDP) %>% head(15)`