

# Session - 6

# Merging with more than one variable

- Most of the merging usage is with two variables: date and company name.

# Merging with more than one variable

- Most of the merging usage is with two variables: date and company name. Generate data using below:

```
date = seq.Date(as.Date("2018-01-01"), by = 1, length.out = 5);
comp = c("A", "B", "C", "D", "E");
all_pairs = merge(comp, date, by = NULL);
# sales data - 15 points
set.seed(1);
idx = sample(1:nrow(all_pairs), 15, replace = F);
df1 = data.frame(comp = all_pairs$x[idx], date = all_pairs$y[idx],
                 sales = round(runif(15, min = 1e3, max = 1e5)));
# advertising data - 12 points
idx = sample(1:nrow(all_pairs), 12, replace = F);
df2 = data.frame(comp = all_pairs$x[idx], date = all_pairs$y[idx],
                 adv = round(runif(12, min = 1e2, max = 1e4)));
```

<u>df1</u>				<u>df2</u>		
<u>comp</u>	<u>date</u>	<u>sales</u>		<u>comp</u>	<u>date</u>	<u>adv</u>
B	05-04-2018	53429		C	03-04-2018	980
A	03-04-2018	70750		C	06-04-2018	6183
B	03-04-2018	5426		B	06-04-2018	8077
B	06-04-2018	88504		B	02-04-2018	241
E	06-04-2018	78449		D	02-04-2018	3642
A	04-04-2018	98954		B	04-04-2018	3878
D	04-04-2018	57618		E	02-04-2018	1501
E	04-04-2018	22011		B	03-04-2018	6727
C	02-04-2018	85976		D	06-04-2018	4259
D	02-04-2018	42842		D	03-04-2018	4434
E	02-04-2018	94057		A	06-04-2018	6745
A	05-04-2018	25063		E	03-04-2018	2208
E	05-04-2018	51320				
D	06-04-2018	99720				
E	03-04-2018	16845				

# Inner Join

- `merge(df1, df2, by = c(" comp", "date"))`

<u>comp</u>	<u>date</u>	<u>sales</u>	<u>adv</u>
B	03-04-2018	5426	6727
B	06-04-2018	88504	8077
D	02-04-2018	42842	3642
E	02-04-2018	94057	1501
D	06-04-2018	99720	4259
E	03-04-2018	16845	2208

# Full Join

- `merge(df1, df2, by = c(" comp", "date"), all = T)`

<u>comp</u>	<u>date</u>	<u>sales</u>	<u>adv</u>
B	05-04-2018	53429	NA
A	03-04-2018	70750	NA
B	03-04-2018	5426	6727
B	06-04-2018	88504	8077
E	06-04-2018	78449	NA
A	04-04-2018	98954	NA
D	04-04-2018	57618	NA
E	04-04-2018	22011	NA
C	02-04-2018	85976	NA
D	02-04-2018	42842	3642
E	02-04-2018	94057	1501
A	05-04-2018	25063	NA
E	05-04-2018	51320	NA
D	06-04-2018	99720	4259
E	03-04-2018	16845	2208
C	03-04-2018	NA	980
C	06-04-2018	NA	6183
B	02-04-2018	NA	241
B	04-04-2018	NA	3878
D	03-04-2018	NA	4434
A	06-04-2018	NA	6745

# Left Join

- `merge(df1, df2, by = c(" comp", "date"), all.x = T)`

<u>comp</u>	<u>date</u>	<u>sales</u>	<u>adv</u>
B	05-04-2018	53429	NA
A	03-04-2018	70750	NA
B	03-04-2018	5426	6727
B	06-04-2018	88504	8077
E	06-04-2018	78449	NA
A	04-04-2018	98954	NA
D	04-04-2018	57618	NA
E	04-04-2018	22011	NA
C	02-04-2018	85976	NA
D	02-04-2018	42842	3642
E	02-04-2018	94057	1501
A	05-04-2018	25063	NA
E	05-04-2018	51320	NA
D	06-04-2018	99720	4259
E	03-04-2018	16845	2208

# Right Join

- `merge(df1, df2, by = c(" comp", "date"), all.y = T)`

<u>comp</u>	<u>date</u>	<u>sales</u>	<u>adv</u>
C	03-04-2018	NA	980
C	06-04-2018	NA	6183
B	06-04-2018	88504	8077
B	02-04-2018	NA	241
D	02-04-2018	42842	3642
B	04-04-2018	NA	3878
E	02-04-2018	94057	1501
B	03-04-2018	5426	6727
D	06-04-2018	99720	4259
D	03-04-2018	NA	4434
A	06-04-2018	NA	6745
E	03-04-2018	16845	2208

# Anti Join

- Sometimes, we wish to find observations in `df1` that are not in `df2`
- `setkey(df1, comp, date); setkey(df2, comp, date);`
- `df1[!df2]` is anti-join.
  - `df2[!df1]` is also anti-join (but the other way round)
- `df1[df2]` is the observations in `df1` that are also in `df2`. This is same as `merge(df1, df2, all.y = T); # right-join`
- `df2[df1]` is left-join. `df1[df2, nomatch = 0]` is inner-join.
- There is no short-hand for full outer join and Cartesian product.
- We can't do anti-join using `merge()` commands.

# Reshaping (long to wide and vice-versa)

- Example data (long form):
  - `dt = fread("long_form_returns.csv", na.strings = "");`
  - Four columns: `cusip`, `Date`, `retadj`, `industry`
- Suppose, we want each firm (`cusip`) in a separate column
  - `wide.ret = dcast(dt, Date + industry ~ cusip, value.var = "retadj");`
    - The wide-form variable comes after `~`
    - Long-form variables comes before `~`
    - Finally, we want “returns” as the value of wide-from format
- How do we get the long form back from `wide.ret`?
  - `dt.org = melt(wide.ret,  
                  id.vars = c("Date", "industry"),  
                  measure.vars = 3:ncol(wide.ret),  
                  variable.name = "cusip", value.name = "retadj",  
                  variable.factor = F, na.rm = T);`



- Wide-form data has only one “value” variable:
  - Rows are dates and columns are company names and the matrix inside is returns. It’s very difficult to get more than one variable.
    - Try: `dt[, ret2 := retadj];`  
`dcast(dt, Date + industry ~ cusip, value.var = c("retadj", "ret2"));`
    - The column names take the form: `retadj_<cusip>` and `ret2_<cusip>`
    - N cusips, D dates and M other variables (like `retadj`, `ret2` etc)
      - long-form: N\*D rows of M variables
      - wide-form: D rows of N\*M variables
- If there are multiple matches in converting to wide-form
  - We can aggregate data using any function like `sum`, `mean`, ...
  - `wide.ret = dcast(dt, Date ~ industry, value.var = "retadj", fun.aggregate = mean, na.rm = T);`
  - This step is irreversible in the sense that we can’t get back original data
- Most of the data will be in long-form but you should know how to quickly convert wide-form to long-form.
  - World bank provides data in wide-form

# Plotting in R

# Plotting in R

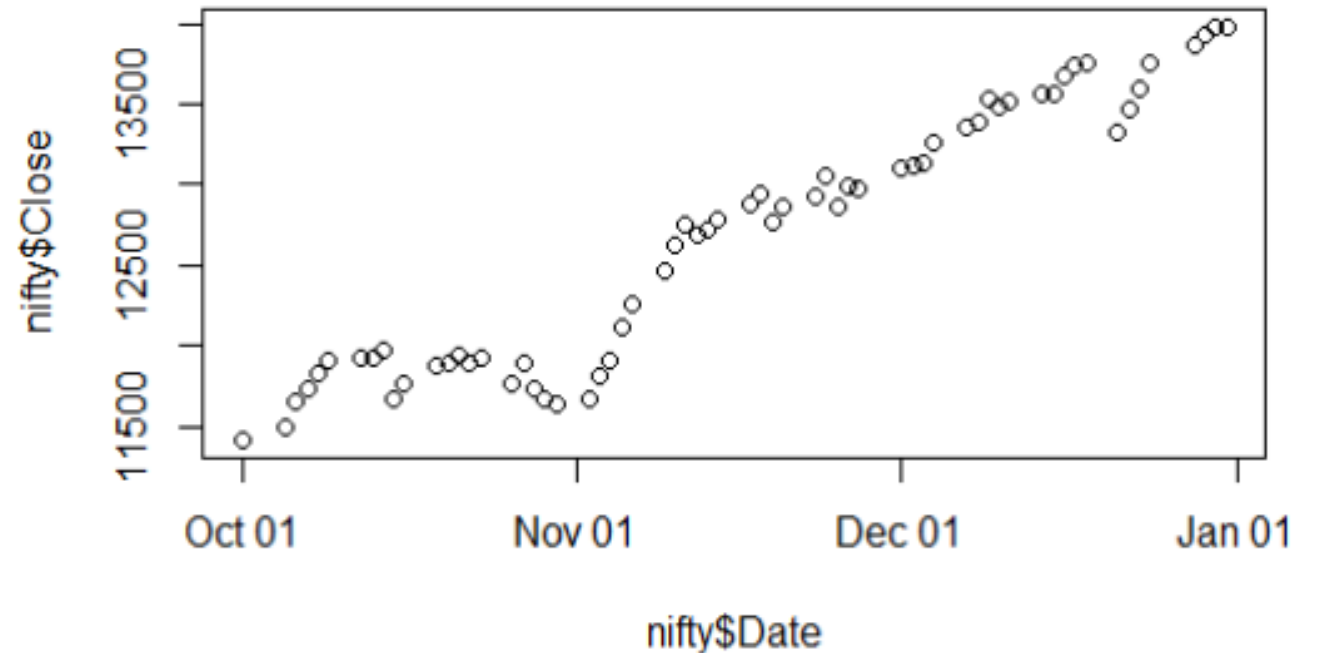
- `plot(x, y, <OPTIONS>);`
- `plot(nifty$Date, nifty$Close);`

# Plotting in R

- `plot(x, y, <OPTIONS>);`
- `plot(nifty$Date, nifty$Close);`
- The below are equivalent to the plot command above:
  - `plot(Date, Close, data = nifty);`
  - `nifty[, plot(Date, Close)];` # only if nifty is a DT


# Plotting in R

- `plot(x, y, <OPTIONS>);`
- `plot(nifty$Date, nifty$Close);`
- The below are equivalent to the plot command above:
  - `plot(Date, Close, data = nifty);`
  - `nifty[, plot(Date, Close)];` # only if nifty is a DT




# Plotting Options

- `plot(nifty$Date, nifty$Close,`  
    `type = 'l',`  
    `col = "blue",`  
    `lty = 2,`  
    `lwd = 2,`  
    `xlab = "Date",`  
    `ylab = "Nifty-50",`  
    `main = "Nifty Plot",`  
    `cex.axis = 2,`  
    `cex.lab = 1.5,`  
    `ylim = c(1e4, 2e4));`




# Multiple Lines

- `plot(nifty$Date, nifty$Close, type = 'l', col = "blue", lty = 1, lwd = 2, xlab = "Date", ylab = "Nifty-50", main = "Nifty Plot");`
- `lines(nifty$Date, nifty$Low, type = 'b', col = "red", lty = 2, lwd = 2);`
- `lines(nifty$Date, nifty$High, type = 'b', col = "green", lty = 2, lwd = 2);`



# Legend

- ```
legend("bottomright", legend = c("Close", "Low", "High"),  
      col = c("blue", "red", "green"),  
      lty = c(1,2,2), lwd = 2,  
      bg = "yellow", text.col = "brown", ncol = 2)
```





# Multiple Axes

```
x = seq(1, 10, length.out = 1e3);
y = sin(x);
z = sqrt(x);

org_mar = par("mar");
par(mar = c(5,5,2,5)) # for extra margin on right y-axis

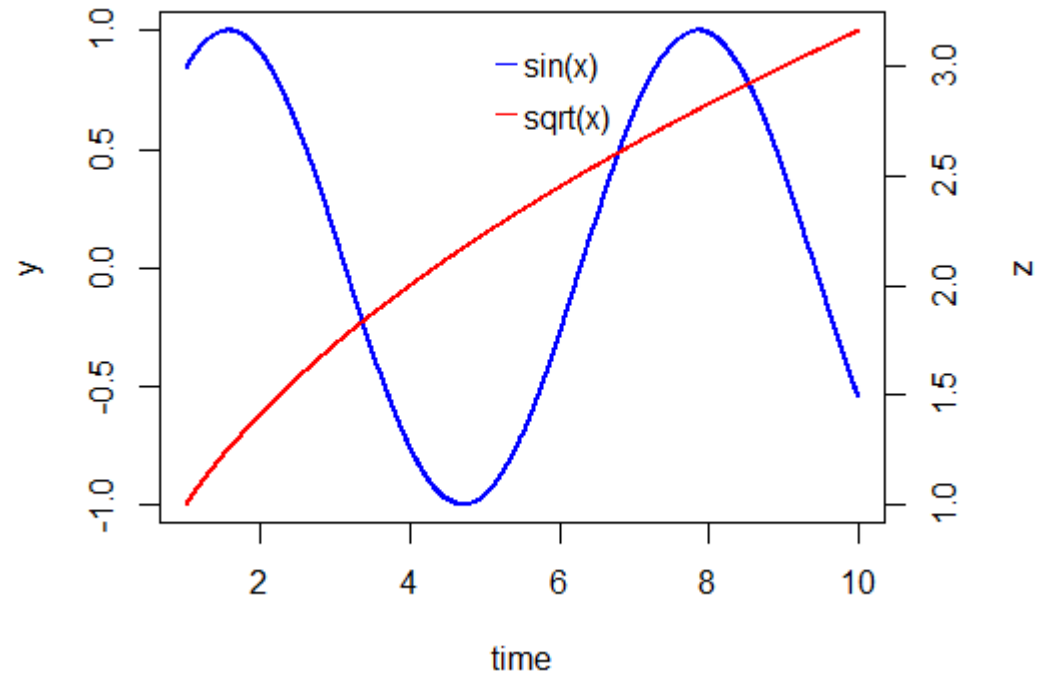
plot(x, y, type = "l", col = "blue", lwd = 2, xlab = "time",
      ylab = "y");
par(new = TRUE);

plot(x, z, type = "l", col = "red", lwd = 2, xlab = NA, ylab =
      NA, axes = F);

axis(4); # makes axis on RHS
mtext(side = 4, line = 3, "z"); # RHS text

legend("top", legend = c("sin(x)", "sqrt(x)"), lty = 1, col =
      c("blue", "red"), lwd = 2, bg = NULL);

par(mar = org_mar);
```



# Multiple Plots



```
x = seq(1, 10, length.out = 1e3);  
y = sin(x);  
z = sqrt(x);
```

```
org_mfrow = par("mfrow");  
par(mfrow = c(2,1));
```

```
plot(x, y, type = "l", col = "blue", lwd = 2, xlab = "time",  
     ylab = "y");  
legend("top", legend = c("sin(x)"), lty = 1, col = c("blue"),  
      lwd = 2);
```

```
plot(x, z, type = "l", col = "red", lwd = 2, xlab = "time",  
     ylab = "z");  
legend("top", legend = c("sqrt(x)"), lty = 1, col = c("red"),  
      lwd = 2);
```

```
par(mfrow = org_mfrow);
```



# Regression Basics

# Regression Basics

- Let,  $Y = \beta_0 + \beta_1 \cdot X + u$  be the true model. By regressing  $Y$  on  $X$ , we hope to recover an unbiased estimate of  $\beta_1$  and see how much of the variation in  $Y$  is explained by variation in  $X$  unrelated to variation in  $u$ .

# Regression Basics


- Let,  $Y = \beta_0 + \beta_1 \cdot X + u$  be the true model. By regressing  $Y$  on  $X$ , we hope to recover an unbiased estimate of  $\beta_1$  and see how much of the variation in  $Y$  is explained by variation in  $X$  unrelated to variation in  $u$ .

```
n = 1000;  
x = rnorm(n, 0, 1);  
u = rnorm(n, 0, 0.1);  
y = u + x;  
plot(x,y);
```

# Regression Basics


- Let,  $Y = \beta_0 + \beta_1 \cdot X + u$  be the true model. By regressing  $Y$  on  $X$ , we hope to recover an unbiased estimate of  $\beta_1$  and see how much of the variation in  $Y$  is explained by variation in  $X$  unrelated to variation in  $u$ .

```
n = 1000;  
x = rnorm(n, 0, 1);  
u = rnorm(n, 0, 0.1);  
y = u + x;  
plot(x,y);
```




```
fit = lm(y ~ x);  
summary(fit);  
stargazer(fit, type = "html", out = "fit.html");  
# Regression Line  
abline(fit$coefficients, col = "red", lwd = 2);
```

```
fit = lm(y ~ x);  
summary(fit);  
stargazer(fit, type = "html", out = "fit.html");  
# Regression Line  
abline(fit$coefficients, col = "red", lwd = 2);
```






```
u = rnorm(n, 0, 0.5);  
y = u + x;  
plot(x,y);  
fit = lm(y ~ x);  
summary(fit);  
# Regression Line  
abline(fit$coefficients, col = "red", lwd = 2);
```



```
u = rnorm(n, 0, 0.5);  
y = u + x;  
plot(x,y);  
fit = lm(y ~ x);  
summary(fit);  
# Regression Line  
abline(fit$coefficients, col = "red", lwd = 2);
```

```
u = rnorm(n, 0, 5);  
y = u + x;  
plot(x,y);  
fit = lm(y ~ x);  
summary(fit);  
# Regression Line  
abline(fit$coefficients, col = "red", lwd = 2);
```

```
u = rnorm(n, 0, 5);  
y = u + x;  
plot(x,y);  
fit = lm(y ~ x);  
summary(fit);  
# Regression Line  
abline(fit$coefficients, col = "red", lwd = 2);
```



# Fixed Effects and Clustering (lfe package)

- This package is currently under repair. So we need to download an earlier version
  - rtools provides a set of tools to build and install earlier version of softwares
  - `remotes::install_version("lfe");`
    - OR try: `remotes::install_github("cran/lfe");`
  - The above will take some time

# Fixed Effects and Clustering (lfe package)

- This package is currently under repair. So we need to download an earlier version
  - rtools provides a set of tools to build and install earlier version of softwares
  - `remotes::install_version("lfe");`
    - OR try: `remotes::install_github("cran/lfe");`
  - The above will take some time
- We ran basic regression as: `lm(y ~ x)`
- Suppose we wish to add fixed effects, we can do
  - `lm(y ~ x + factor(fe_1) + factor(fe_2) + ...)`
  - The above works fine for a small number of fixed effects (FE)
  - `lm()` is prohibitively slow for large number of FE and FE with large num of levels

- There are two ways to take care of FE
  - Include them in regression
  - De-mean the variables (both x and y) with respect to those FE levels
    - `lfe:felm()` is very efficient at this

- There are two ways to take care of FE
  - Include them in regression
  - De-mean the variables (both x and y) with respect to those FE levels
    - `lfe:felm()` is very efficient at this
- `lm()` offers no support for clustering of standard errors
  - Most recent research includes some form of clustering in the results
    - In panel data, you often need multi-way clustering
    - Earlier approach was to correct for heteroscedasticity and autocorrelation (HAC)
      - You might recognize terms like White (Robust) standard errors, Newey-West adjustment, etc while reading papers
  - `lfe::felm()` provides inbuilt cluster robust standard errors



- There are two ways to take care of FE
  - Include them in regression
  - De-mean the variables (both x and y) with respect to those FE levels
    - `lfe:felm()` is very efficient at this
- `lm()` offers no support for clustering of standard errors
  - Most recent research includes some form of clustering in the results
    - In panel data, you often need multi-way clustering
    - Earlier approach was to correct for heteroscedasticity and autocorrelation (HAC)
      - You might recognize terms like White (Robust) standard errors, Newey-West adjustment, etc while reading papers
  - `lfe::felm()` provides inbuilt cluster robust standard errors
- Syntax: `felm(formula, data)`
  - formula:      <MODEL>            |      <FE>            | <INSTR> |      <CLUSTERS>
  - `felm( y ~ x1 + x2       | f1 + f2   |    0       |      c1 + c2       )`

?felm (Help page of felm command)

# ?felm (Help page of felm command)

- The formula specification is a response variable followed by a four part formula.
  - The first part consists of ordinary covariates, the second part consists of factors to be projected out. The third part is an IV-specification. The fourth part is a cluster specification for the standard errors.
  - I.e. something like  $y \sim x1 + x2 \mid f1 + f2 \mid (Q \mid W \sim x3 + x4) \mid clu1 + clu2$  where  $y$  is the response,  $x1, x2$  are ordinary covariates,  $f1, f2$  are factors to be projected out,  $Q$  and  $W$  are covariates which are instrumented by  $x3$  and  $x4$ , and  $clu1, clu2$  are factors to be used for computing cluster robust standard errors.
  - Parts that are not used should be specified as  $\emptyset$ , except if it's at the end of the formula, where they can be omitted.
    - The parentheses are needed in the third part since  $\mid$  has higher precedence than  $\sim$ .
  - Multiple left hand sides like  $y \mid w \mid x \sim x1 + x2 \mid f1 + f2 \mid \dots$  are allowed.

# Different Regressions in R

# Different Regressions in R

- OLS is very simple
  - `fit = lm(y ~ x)`
    - You can check the summary using `summary(fit)`. This gives std errors, t-stats and R2
  - `lfe::fe1m(y ~ x)`

# Different Regressions in R

- OLS is very simple

- `fit = lm(y ~ x)`

- You can check the summary using `summary(fit)`. This gives std errors, t-stats and R2

- `lfe::fe1m(y ~ x)`

- GLS

- `nlme::gls(y ~ x, correlation = C, weights = w)`

- C is the group correlation matrix, while w are heteroscedasticity weights

# Different Regressions in R

- OLS is very simple
  - `fit = lm(y ~ x)`
    - You can check the summary using `summary(fit)`. This gives std errors, t-stats and R2
  - `lfe::felm(y ~ x)`
- GLS
  - `nlme::gls(y ~ x, correlation = C, weights = w)`
    - C is the group correlation matrix, while w are heteroscedasticity weights
- IV
  - We can use `lfe::felm` for IV regression
    - `lfe::felm(y ~ x1 + x2 | 0 | (x3|x4 ~ w3 + w4) | 0)`
    - `AER::ivreg(y ~ x1 + x2 + x3 + x4 | x1 + x2 + w3 + w4)`
  - Requirements:
    - Instruments (w3, w4) must be correlated with endogenous variables (x3, x4)
    - Instruments (w3, w4) are unrelated to the error term, i.e. instruments affect y only through endogenous variables x3 and x4

- Logit and Probit Models

- `glm(y ~ x1 + x2, family = binomial("logit"))`
- `glm(y ~ x1 + x2, family = binomial("probit"))`
- Try `?family` for more insights into how to use other models



- Logit and Probit Models

- `glm(y ~ x1 + x2, family = binomial("logit"))`
- `glm(y ~ x1 + x2, family = binomial("probit"))`
- Try `?family` for more insights into how to use other models

- For more details check out `VGAM::vglm` which implements vectorised glm for several other models

- The need to perform a regression other than plain vanilla OLS rarely arises

- Important exceptions are IV (2-SLS), FE and clustering.
- Refer “Introductory Econometrics” by Woolridge for more theory and details on different regression models.

- Logit and Probit Models

- `glm(y ~ x1 + x2, family = binomial("logit"))`
- `glm(y ~ x1 + x2, family = binomial("probit"))`
- Try `?family` for more insights into how to use other models

- For more details check out `VGAM::vglm` which implements vectorised glm for several other models

- The need to perform a regression other than plain vanilla OLS rarely arises

- Important exceptions are IV (2-SLS), FE and clustering.
- Refer “Introductory Econometrics” by Woolridge for more theory and details on different regression models.

- Bootstrapping

- In some models, its impossible to accurately judge the structure of standard errors. There we can employ boot-strapping to get standard errors.
  - Sample N data points (with repetition) from your dataset and estimate the model
    - Do this 1000 (or more) times
    - The standard error of coef. errors is then your standard errors

# Time-series Regressions

- The time-series regression is specified the same way as a cross-sectional regression

# Time-series Regressions

- The time-series regression is specified the same way as a cross-sectional regression
- But we need to be careful about some issues
  - Non-stationarity in data
    - 1<sup>st</sup>/2<sup>nd</sup> order integration?
  - Persistent time-series
    - AR (autoregressive) and moving average (MA) parameters
  - Fit `arima(x, order)` model where `order = c(p, d, q)`
    - `p`: AR order, `d`: integration order, `q`: MA order
  - Spurious Regression
    - Will return of SBI predict return of HDFC? Reliance?
      - Nifty (or broader market) return predicts both SBI and Reliance!
  - Co-integration
    - Does India's GDP forecasts predict Nifty levels?
- Learn more at [https://en.wikipedia.org/wiki/Autoregressive%E2%80%93moving-average\\_model](https://en.wikipedia.org/wiki/Autoregressive%E2%80%93moving-average_model)